

Database System

Concurrency Control

Muhammad Tariq Mahmood

tariq@koreatech.ac.kr

School of Computer Science and Engineering
Korea University of Technology and Education

Database Concurrency Control

- ▶ Purpose of Concurrency Control
 - To enforce Isolation (through mutual exclusion) among conflicting transactions.
 - To preserve database consistency through consistency preserving execution of transactions.
 - To resolve read–write and write–write conflicts.
- ▶ Example:
 - In concurrent execution environment if T1 conflicts with T2 over a data item A, then the existing concurrency control decides if T1 or T2 should get the A and if the other transaction is rolled–back or waits.

Concurrency Control Algorithms

- Lock based CC Algorithms
 - Two-phase locking Algorithm
 - Multiversions CC algorithms
 - Validation CC Algorithms
 -
- Timestamp based CC Algorithms
 - Basic Timestamp based Algorithms
 -

Two-Phase Locking Techniques

- ▶ Locking is an operation which secures
 - (a) permission to Read
 - (b) permission to Write a data item for a transaction.
- ▶ Example:
 - Lock (X). Data item X is locked in behalf of the requesting transaction.
- ▶ Unlocking is an operation which removes these permissions from the data item.
- ▶ Example:
 - Unlock (X): Data item X is made available to all other transactions.
- ▶ Lock and Unlock are Atomic operations.

Two-Phase Locking Techniques:

Essential components

- ▶ Two locks modes:

- (a) shared (read)
- (b) exclusive (write).

- ▶ Shared mode: shared lock (X)

- More than one transaction can apply share lock on X for reading its value but no write lock can be applied on X by any other transaction.

- ▶ Exclusive mode: Write lock (X)

- Only one write lock on X can exist at any time and no shared lock can be applied by any other transaction on X.

- ▶ Conflict matrix

	Read	Write
Read	Y	N
Write	N	N

Two-Phase Locking Techniques:

Essential components

- ▶ Lock Manager:
 - Managing locks on data items.
- ▶ Lock table:
 - Lock manager uses it to store the identify of transaction locking a data item, the data item, lock mode and pointer to the next data item locked. One simple way to implement a lock table is through linked list.

Transaction ID	Data item id	lock mode	Ptr to next data item
T1	X1	Read	Next

Two-Phase Locking Techniques:

Essential components

- ▶ Database requires that all transactions should be well-formed. A transaction is well-formed if:
 - It must lock the data item before it reads or writes to it.
 - It must not lock an already locked data items and it must not try to unlock a free data item.

lock_item(X):

```
B:  if LOCK(X) = 0          (* item is unlocked *)
      then LOCK(X) ← 1      (* lock the item *)
      else
        begin
          wait (until LOCK(X) = 0
                and the lock manager wakes up the transaction);
          go to B
        end;
```

unlock_item(X):

```
LOCK(X) ← 0;          (* unlock the item *)
if any transactions are waiting
  then wakeup one of the waiting transactions;
```

Figure 22.1

Lock and unlock operations for binary locks.

```

read_lock(X):
B:  if LOCK(X) = "unlocked"
      then begin LOCK(X)  $\leftarrow$  "read-locked";
           no_of_reads(X)  $\leftarrow$  1
           end
      else if LOCK(X) = "read-locked"
           then no_of_reads(X)  $\leftarrow$  no_of_reads(X) + 1
      else begin
           wait (until LOCK(X) = "unlocked"
                and the lock manager wakes up the transaction);
           go to B
           end;

write_lock(X):
B:  if LOCK(X) = "unlocked"
      then LOCK(X)  $\leftarrow$  "write-locked"
      else begin
           wait (until LOCK(X) = "unlocked"
                and the lock manager wakes up the transaction);
           go to B
           end;

unlock (X):
      if LOCK(X) = "write-locked"
          then begin LOCK(X)  $\leftarrow$  "unlocked";
               wakeup one of the waiting transactions, if any
               end
      else if LOCK(X) = "read-locked"
          then begin
               no_of_reads(X)  $\leftarrow$  no_of_reads(X) - 1;
               if no_of_reads(X) = 0
                   then begin LOCK(X) = "unlocked";
                        wakeup one of the waiting transactions, if any
                        end
               end;

```

Figure 22.2
 Locking and unlocking
 operations for two-
 mode (read-write or
 shared-exclusive)
 locks.

Two-Phase Locking Techniques:

The algorithm

- ▶ If the simple binary locking scheme described here is used, every transaction must obey the following rules:
 1. A transaction T must issue the operation $\text{lock_item}(X)$ before any $\text{read_item}(X)$ or $\text{write_item}(X)$ operations are performed in T .
 2. A transaction T must issue the operation $\text{unlock_item}(X)$ after all $\text{read_item}(X)$ and $\text{write_item}(X)$ operations are completed in T .
 3. A transaction T will not issue a $\text{lock_item}(X)$ operation if it already holds the lock on item X .¹
 4. A transaction T will not issue an $\text{unlock_item}(X)$ operation unless it already holds the lock on item X .

Two-Phase Locking Techniques:

The algorithm

- ▶ When we use the shared/exclusive locking scheme, the system must enforce the following rules:
 1. A transaction T must issue the operation $\text{read_lock}(X)$ or $\text{write_lock}(X)$ before any $\text{read_item}(X)$ operation is performed in T .
 2. A transaction T must issue the operation $\text{write_lock}(X)$ before any $\text{write_item}(X)$ operation is performed in T .
 3. A transaction T must issue the operation $\text{unlock}(X)$ after all $\text{read_item}(X)$ and $\text{write_item}(X)$ operations are completed in T .
 4. A transaction T will not issue a $\text{read_lock}(X)$ operation if it already holds a read (shared) lock or a write (exclusive) lock on item X . This rule may be relaxed, as we discuss shortly.
 5. A transaction T will not issue a $\text{write_lock}(X)$ operation if it already holds a read (shared) lock or write (exclusive) lock on item X . This rule may also be relaxed, as we discuss shortly.
 6. A transaction T will not issue an $\text{unlock}(X)$ operation unless it already holds a read (shared) lock or a write (exclusive) lock on item X .

Two-Phase Locking Techniques: The algorithm

- ▶ Two Phases:
 - (a) Locking (Growing)
 - (b) Unlocking (Shrinking).
- ▶ Locking (Growing) Phase:
 - A transaction applies locks (read or write) on desired data items one at a time.
- ▶ Unlocking (Shrinking) Phase:
 - A transaction unlocks its locked data items one at a time.
- ▶ Requirement:
 - For a transaction these two phases must be mutually exclusively, that is, during locking phase unlocking phase must not start and during unlocking phase locking phase must not begin.

Two-Phase Locking Techniques:

The algorithm

- Transactions T_1 and T_2 in Figure do not follow the two-phase locking protocol because the `write_lock(X)` operation follows the `unlock(Y)` operation in T_1 , and similarly the `write_lock(Y)` operation follows the `unlock(X)` operation in T_2 .

(a)

T_1	T_2
<code>read_lock(Y);</code> <code>read_item(Y);</code> <code>unlock(Y);</code> <code>write_lock(X);</code> <code>read_item(X);</code> <code>X := X + Y;</code> <code>write_item(X);</code> <code>unlock(X);</code>	<code>read_lock(X);</code> <code>read_item(X);</code> <code>unlock(X);</code> <code>write_lock(Y);</code> <code>read_item(Y);</code> <code>Y := X + Y;</code> <code>write_item(Y);</code> <code>unlock(Y);</code>

(b)

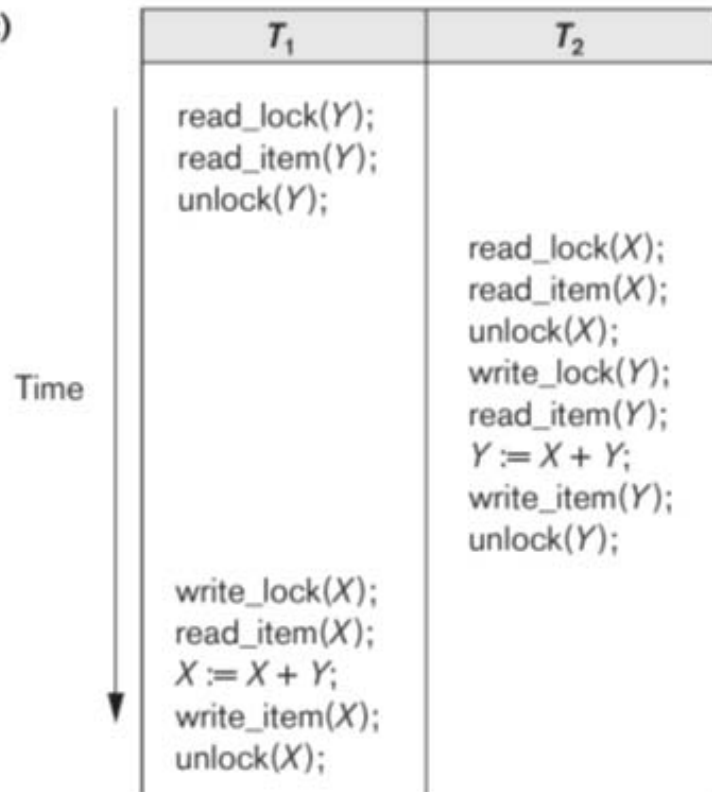
Initial values: $X=20, Y=30$

Result serial schedule T_1
followed by T_2 : $X=50, Y=80$

Result of serial schedule T_2
followed by T_1 : $X=70, Y=50$

Two-Phase Locking Techniques: The algorithm

(c)



Result of schedule S:
 $X=50, Y=50$
(nonserializable)

Figure 22.3

Transactions that do not obey two-phase locking. (a) Two transactions T_1 and T_2 . (b) Results of possible serial schedules of T_1 and T_2 . (c) A nonserializable schedule S that uses locks.

Two-Phase Locking Techniques:

The algorithm

- ▶ If we enforce two-phase locking, the transactions can be rewritten as T_1' and T_2' , as shown in Figure
- ▶ Now, the schedule shown in Figure 22.3(c) is not permitted for T_1' and T_2' (with their modified order of locking and unlocking operations) under the rules of locking

T_1'	T_2'
read_lock(Y); read_item(Y); write_lock(X); unlock(Y) read_item(X); $X := X + Y$; write_item(X); unlock(X);	read_lock(X); read_item(X); write_lock(Y); unlock(X) read_item(Y); $Y := X + Y$; write_item(Y); unlock(Y);

Figure 22.4

Transactions T_1' and T_2' , which are the same as T_1 and T_2 in Figure 22.3, but follow the two-phase locking protocol. Note that they can produce a deadlock.

Two-Phase Locking Techniques: The algorithm

- ▶ Two-phase policy generates two locking algorithms
 - (a) **Basic**
 - (b) **Conservative**
- ▶ **Conservative:**
 - Prevents deadlock by locking all desired data items before transaction begins execution.
- ▶ **Basic:**
 - Transaction locks data items incrementally. This may cause deadlock which is dealt with.
- ▶ **Strict:**
 - A more stricter version of Basic algorithm where unlocking is performed after a transaction terminates (commits or aborts and rolled-back). This is the most commonly used two-phase locking algorithm.

Two-Phase Locking Techniques: The algorithm

- ▶ Two-phase policy generates two locking algorithms
 - (a) **Basic**
 - (b) **Conservative**
- ▶ **Conservative:**
 - Prevents deadlock by locking all desired data items before transaction begins execution.
- ▶ **Basic:**
 - Transaction locks data items incrementally. This may cause deadlock which is dealt with.
- ▶ **Strict:**
 - A more stricter version of Basic algorithm where unlocking is performed after a transaction terminates (commits or aborts and rolled-back). This is the most commonly used two-phase locking algorithm.

Two-Phase Locking Techniques: Dealing with Deadlock and Starvation

- ▶ Deadlock occurs when *each* transaction T in a set of *two or more transactions* is waiting for some item that is locked by some other transaction T' in the set.
- ▶ A simple example is shown in Figure where the two transactions T_1' and T_2' are deadlocked in a partial schedule;
- ▶ T_1' is in the waiting queue for X , which is locked by T_2' , while T_2' is in the waiting queue for Y , which is locked by T_1' .

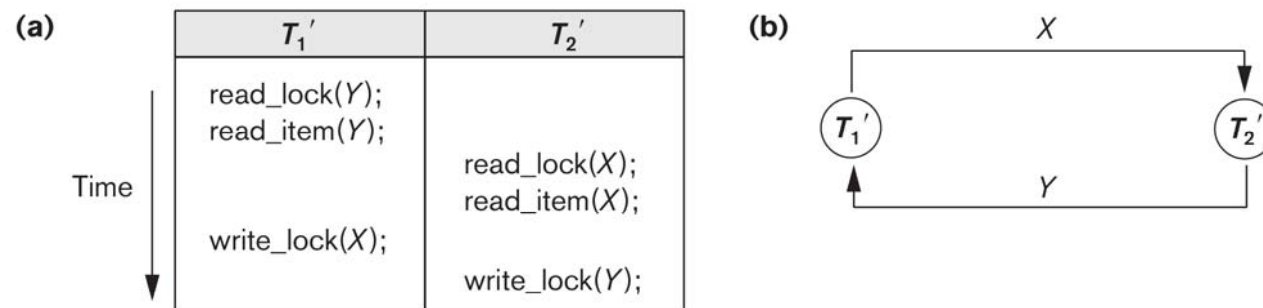


Figure 22.5

Illustrating the deadlock problem. (a) A partial schedule of T_1' and T_2' that is in a state of deadlock. (b) A wait-for graph for the partial schedule in (a).

Two-Phase Locking Techniques: Dealing with Deadlock and Starvation

- ▶ **Deadlock prevention**

- A transaction locks all data items it refers to before it begins execution.
- This way of locking prevents deadlock since a transaction never waits for a data item.
- The conservative two-phase locking uses this approach.

Two-Phase Locking Techniques: Dealing with Deadlock and Starvation

▶ Deadlock detection and resolution

- In this approach, deadlocks are allowed to happen. The scheduler maintains a wait-for-graph for detecting cycle. If a cycle exists, then one transaction involved in the cycle is selected (victim) and rolled-back.
- A wait-for-graph is created using the lock table. As soon as a transaction is blocked, it is added to the graph. When a chain like: T_i waits for T_j waits for T_k waits for T_i or T_j occurs, then this creates a cycle. One of the transaction o

Two-Phase Locking Techniques: Dealing with Deadlock and Starvation

► Deadlock avoidance

- There are many variations of two-phase locking algorithm.
- Some avoid deadlock by not letting the cycle to complete.
- That is as soon as the algorithm discovers that blocking a transaction is likely to create a cycle, it rolls back the transaction.
- Wound-Wait and Wait-Die algorithms use timestamps to avoid deadlocks by rolling-back victim.

Two-Phase Locking Techniques: Dealing with Deadlock and Starvation

▶ Starvation

- Starvation occurs when a particular transaction consistently waits or restarted and never gets a chance to proceed further.
- In a deadlock resolution it is possible that the same transaction may consistently be selected as victim and rolled-back.
- This limitation is inherent in all priority based scheduling mechanisms.
- In Wound-Wait scheme a younger transaction may always be wounded (aborted) by a long running older transaction which may create starvation.

Two-Phase Locking Techniques: Dealing with Deadlock and Starvation

▶ Timestamp

- A monotonically increasing variable (integer) indicating the age of an operation or a transaction.
- A larger timestamp value indicates a more recent event or operation.
- Timestamp based algorithm uses timestamp to serialize the execution of concurrent transactions.